

# Data cleaning

```
In [1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
%matplotlib inline
import matplotlib
matplotlib.rcParams["figure.figsize"] = (20,10)
```

## Importing the Data Set

```
In [2]: df= pd.read_csv('datasets_20710_26737_Bengaluru_House_Data.csv')
df.head()
```

```
Out[2]:
```

	area_type	availability	location	size	society	total_sqft	bath	balcony	price
0	Super built-up Area	19-Dec	Electronic City Phase II	2 BHK	Coomee	1056	2.0	1.0	39.07
1	Plot Area	Ready To Move	Chikka Tirupathi	4 Bedroom	Theanmp	2600	5.0	3.0	120.00
2	Built-up Area	Ready To Move	Uttarahalli	3 BHK	NaN	1440	2.0	3.0	62.00
3	Super built-up Area	Ready To Move	Lingadheeranahalli	3 BHK	Soiewre	1521	3.0	1.0	95.00
4	Super built-up Area	Ready To Move	Kothanur	2 BHK	NaN	1200	2.0	1.0	51.00

```
In [3]: df.shape #size of the data set to get to know how may data point we are having
```

```
Out[3]: (13320, 9)
```

```
In [4]: df.groupby('area_type')['area_type'].agg('count')
#so we are grouping the data
#This a catogorical variable as we can see they are divided into various catogories
and the respective count is displayed
```

```
Out[4]: area_type
Built-up Area      2418
Carpet Area         87
Plot Area          2025
Super built-up Area 8790
Name: area_type, dtype: int64
```

Now Analyse what variables will be dependent on the dependent variable X i.e price

```
In [5]: df1= df.drop(['availability','area_type','society','balcony'], axis='columns')
df1.head()
#so here we have removed a few variables to make our code easy but we can still depend on few variables
#but they are being dropped here
```

Out[5]:

	location	size	total_sqft	bath	price
0	Electronic City Phase II	2 BHK	1056	2.0	39.07
1	Chikka Tirupathi	4 Bedroom	2600	5.0	120.00
2	Uttarahalli	3 BHK	1440	2.0	62.00
3	Lingadheeranahalli	3 BHK	1521	3.0	95.00
4	Kothanur	2 BHK	1200	2.0	51.00

```
In [6]: df1.isnull().sum() # this shows the sum of the null values in the data set
```

```
Out[6]: location      1
size      16
total_sqft  0
bath      73
price      0
dtype: int64
```

```
In [7]: df2=df1.dropna()
# here we are dropping the null values
```

```
In [8]: df2.isnull().sum()
```

```
Out[8]: location      0
size      0
total_sqft  0
bath      0
price      0
dtype: int64
```

```
In [9]: df3 = df2
```

```
In [10]: df3['size'].unique() #this column is useless unless and until we take the integer values and categorise them properly
```

```
Out[10]: array(['2 BHK', '4 Bedroom', '3 BHK', '4 BHK', '6 Bedroom', '3 Bedroom',
                '1 BHK', '1 RK', '1 Bedroom', '8 Bedroom', '2 Bedroom',
                '7 Bedroom', '5 BHK', '7 BHK', '6 BHK', '5 Bedroom', '11 BHK',
                '9 BHK', '9 Bedroom', '27 BHK', '10 Bedroom', '11 Bedroom',
                '10 BHK', '19 BHK', '16 BHK', '43 Bedroom', '14 BHK', '8 BHK',
                '12 Bedroom', '13 BHK', '18 Bedroom'], dtype=object)
```

```
In [11]: df3['Bhk']= df3['size'].apply(lambda x: int(x.split(' ')[0]))
# here we are acquiring the '4' from '4 bhk' which is divided by a space i.e ' ' and then placing them in a column and also converting them into integers
```

D:\New folder\lib\site-packages\ipykernel\_launcher.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
 """Entry point for launching an IPython kernel.

```
In [12]: df3['Bhk'].unique()# now we can see we have a unique column of no. of bedrooms
```

```
Out[12]: array([ 2,  4,  3,  6,  1,  8,  7,  5, 11,  9, 27, 10, 19, 16, 43, 14, 12,
        13, 18], dtype=int64)
```

```
In [13]: df3.head()
```

```
Out[13]:
```

	location	size	total_sqft	bath	price	Bhk
0	Electronic City Phase II	2 BHK	1056	2.0	39.07	2
1	Chikka Tirupathi	4 Bedroom	2600	5.0	120.00	4
2	Uttarahalli	3 BHK	1440	2.0	62.00	3
3	Lingadheeranahalli	3 BHK	1521	3.0	95.00	3
4	Kothanur	2 BHK	1200	2.0	51.00	2

```
In [14]: df3['total_sqft'].unique()
```

```
Out[14]: array(['1056', '2600', '1440', ..., '1133 - 1384', '774', '4689'],
        dtype=object)
```

```
In [15]: def is_float(x):
        try:
            float(x)
        except:
            return False
        return True
```

```
In [16]: df3[~df3['total_sqft'].apply(is_float)]
```

```
Out[16]:
```

	location	size	total_sqft	bath	price	Bhk
30	Yelahanka	4 BHK	2100 - 2850	4.0	186.000	4
122	Hebbal	4 BHK	3067 - 8156	4.0	477.000	4
137	8th Phase JP Nagar	2 BHK	1042 - 1105	2.0	54.005	2
165	Sarjapur	2 BHK	1145 - 1340	2.0	43.490	2
188	KR Puram	2 BHK	1015 - 1540	2.0	56.800	2
...	...	...	...	...	...	...
12975	Whitefield	2 BHK	850 - 1060	2.0	38.190	2
12990	Talaghattapura	3 BHK	1804 - 2273	3.0	122.000	3
13059	Harlur	2 BHK	1200 - 1470	2.0	72.760	2
13265	Hoodi	2 BHK	1133 - 1384	2.0	59.135	2
13299	Whitefield	4 BHK	2830 - 2882	5.0	154.500	4

190 rows × 6 columns

```
In [17]: def convert_sqft_to_num(x):
        tokens=x.split('-')
        if len(tokens) == 2:
            return (float(tokens[0])+float(tokens[1]))/2
        try:
            return float(x)
        except:
            return None
```

```
In [18]: convert_sqft_to_num('2166')
```

```
Out[18]: 2166.0
```

```
In [19]: convert_sqft_to_num('1145 - 1340')
```

```
Out[19]: 1242.5
```

```
In [20]: df4 = df3.copy()
df4['total_sqft'] = df4['total_sqft'].apply(convert_sqft_to_num)
```

```
In [21]: df4.head(5)
```

```
Out[21]:
```

	location	size	total_sqft	bath	price	Bhk
0	Electronic City Phase II	2 BHK	1056.0	2.0	39.07	2
1	Chikka Tirupathi	4 Bedroom	2600.0	5.0	120.00	4
2	Uttarahalli	3 BHK	1440.0	2.0	62.00	3
3	Lingadheeranahalli	3 BHK	1521.0	3.0	95.00	3
4	Kothanur	2 BHK	1200.0	2.0	51.00	2

```
In [22]: df4.loc[30]
```

```
Out[22]: location      Yelahanka
size                4 BHK
total_sqft          2475
bath                4
price              186
Bhk                4
Name: 30, dtype: object
```

## Feature Engineering

```
In [23]: df5 = df4.copy()
df5['price_per_sqft'] = df5['price']*100000/df5['total_sqft']
df5.head()
```

```
Out[23]:
```

	location	size	total_sqft	bath	price	Bhk	price_per_sqft
0	Electronic City Phase II	2 BHK	1056.0	2.0	39.07	2	3699.810606
1	Chikka Tirupathi	4 Bedroom	2600.0	5.0	120.00	4	4615.384615
2	Uttarahalli	3 BHK	1440.0	2.0	62.00	3	4305.555556
3	Lingadheeranahalli	3 BHK	1521.0	3.0	95.00	3	6245.890861
4	Kothanur	2 BHK	1200.0	2.0	51.00	2	4250.000000

```
In [24]: df5.location.unique()
```

```
Out[24]: array(['Electronic City Phase II', 'Chikka Tirupathi', 'Uttarahalli', ...,
        '12th cross srinivas nagar banshankari 3rd stage',
        'Havanur extension', 'Abshot Layout'], dtype=object)
```

```
In [25]: len(df5.location.unique())
```

```
Out[25]: 1304
```

```
In [26]: df5['location'] = df5['location'].apply( lambda x: x.strip() )
location_stats = df5.groupby('location')['location'].agg('count').sort_values (ascending = False)
location_stats
```

```
Out[26]: location
Whitefield          535
Sarjapur Road       392
Electronic City     304
Kanakpura Road      266
Thanisandra         236
...
LIC Colony          1
Kuvempu Layout      1
Kumbhena Agrahara   1
Kudlu Village,      1
1 Annasandrapalya   1
Name: location, Length: 1293, dtype: int64
```

```
In [27]: len(location_stats[location_stats<=10])
```

```
Out[27]: 1052
```

```
In [28]: location_stats_less_than_10 = location_stats[location_stats<=10]
location_stats_less_than_10
```

```
Out[28]: location
BTM 1st Stage       10
Basapura            10
Sector 1 HSR Layout 10
Naganathapura       10
Kalkere             10
..
LIC Colony          1
Kuvempu Layout      1
Kumbhena Agrahara   1
Kudlu Village,      1
1 Annasandrapalya   1
Name: location, Length: 1052, dtype: int64
```

```
In [29]: len(df5.location.unique())
```

```
Out[29]: 1293
```

```
In [30]: df5.location = df5.location.apply(lambda x: 'other' if x in location_stats_less_than_10 else x)
len(df5.location.unique())
```

```
Out[30]: 242
```

```
In [31]: df5.head(10)
```

```
Out[31]:
```

	location	size	total_sqft	bath	price	Bhk	price_per_sqft
0	Electronic City Phase II	2 BHK	1056.0	2.0	39.07	2	3699.810606
1	Chikka Tirupathi	4 Bedroom	2600.0	5.0	120.00	4	4615.384615
2	Uttarahalli	3 BHK	1440.0	2.0	62.00	3	4305.555556
3	Lingadheeranahalli	3 BHK	1521.0	3.0	95.00	3	6245.890861
4	Kothanur	2 BHK	1200.0	2.0	51.00	2	4250.000000
5	Whitefield	2 BHK	1170.0	2.0	38.00	2	3247.863248
6	Old Airport Road	4 BHK	2732.0	4.0	204.00	4	7467.057101
7	Rajaji Nagar	4 BHK	3300.0	4.0	600.00	4	18181.818182
8	Marathahalli	3 BHK	1310.0	3.0	63.25	3	4828.244275
9	other	6 Bedroom	1020.0	6.0	370.00	6	36274.509804

```
In [32]: df5[df5.total_sqft/df5.Bhk > 300].head()  
df5.shape
```

```
Out[32]: (13246, 7)
```

```
In [33]: df6 = df5[~(df5.total_sqft/df5.Bhk < 300)]  
df6.shape
```

```
Out[33]: (12502, 7)
```

```
In [34]: df6.price_per_sqft.describe()
```


```
Out[34]: count      12456.000000  
mean         6308.502826  
std          4168.127339  
min           267.829813  
25%          4210.526316  
50%          5294.117647  
75%          6916.666667  
max         176470.588235  
Name: price_per_sqft, dtype: float64
```

```
In [35]: def remove_pps_outliers(df):  
    df_out = pd.DataFrame()  
    for key, subdf in df.groupby('location'):  
        m = np.mean(subdf.price_per_sqft)  
        st = np.std(subdf.price_per_sqft)  
        reduced_df = subdf[(subdf.price_per_sqft>(m-st)) & (subdf.price_per_sqft<(m+st))]  
        df_out = pd.concat([df_out,reduced_df],ignore_index = True)  
    return df_out
```

```
In [36]: df7 = remove_pps_outliers(df6)  
df7.shape
```

```
Out[36]: (10241, 7)
```


```
In [37]: def plot_scatter_chart(df,location):
    bhk2 = df[(df.location==location) & (df.Bhk==2)]
    bhk3 = df[(df.location==location) & (df.Bhk==3)]
    matplotlib.rcParams['figure.figsize'] = (15,10)
    plt.scatter(bhk2.total_sqft,bhk2.price,color='blue' ,label='2 BHK',s=50)
    plt.scatter(bhk3.total_sqft,bhk3.price,marker= '+',color='green' ,label='3 BHK',s
=50)
    plt.xlabel("Total Square Feet Area")
    plt.ylabel("Price")
    plt.title(location)
    plt.legend()
plot_scatter_chart(df7,"Hebbal")
```



```
In [38]: def remove_bhk_outliners(df):
    exclude_indices = np.array([])
    for location,location_df in df.groupby('location'):
        bhk_stats = {}
        for Bhk, bhk_df in location_df.groupby('Bhk'):
            bhk_stats[Bhk] = {
                'mean': np.mean(bhk_df.price_per_sqft),
                'std': np.std(bhk_df.price_per_sqft),
                'count': bhk_df.shape[0]
            }
        for Bhk, bhk_df in location_df.groupby('Bhk'):
            stats = bhk_stats.get(Bhk-1)
            if stats and stats['count']> 5:
                exclude_indices = np.append(exclude_indices,bhk_df[bhk_df.price_per_s
qft<(stats['mean'])].index.values)
    return df.drop(exclude_indices,axis='index')
df8 = remove_bhk_outliners(df7)
df8.shape
```


Out[38]: (7329, 7)

```
In [39]: plot_scatter_chart(df8,"Hebbal")
```



```
In [40]: import matplotlib
matplotlib.rcParams["figure.figsize"] = (20,10)
plt.hist(df8.price_per_sqft,rwidth=0.8)
plt.xlabel("Price Per Square Feet")
plt.ylabel("Count")
```

```
Out[40]: Text(0, 0.5, 'Count')
```



```
In [41]: df8.bath.unique()
```

```
Out[41]: array([ 4.,  3.,  2.,  5.,  8.,  1.,  6.,  7.,  9., 12., 16., 13.]
```




```
In [42]: df8[df8.bath>10]
```

```
Out[42]:
```

	location	size	total_sqft	bath	price	Bhk	price_per_sqft
5277	Neeladri Nagar	10 BHK	4000.0	12.0	160.0	10	4000.000000
8486	other	10 BHK	12000.0	12.0	525.0	10	4375.000000
8575	other	16 BHK	10000.0	16.0	550.0	16	5500.000000
9308	other	11 BHK	6000.0	12.0	150.0	11	2500.000000
9639	other	13 BHK	5425.0	13.0	275.0	13	5069.124424

```
In [43]: import matplotlib
matplotlib.rcParams["figure.figsize"]= (20,10)
plt.hist(df8.bath,rwidth=0.8)
plt.xlabel('number of bathrooms ')
plt.ylabel('count')
```

```
Out[43]: Text(0, 0.5, 'count')
```



```
In [44]: df8[df8.bath>df8.Bhk+2]
```

```
Out[44]:
```

	location	size	total_sqft	bath	price	Bhk	price_per_sqft
1626	Chikkabanavar	4 Bedroom	2460.0	7.0	80.0	4	3252.032520
5238	Nagasandra	4 Bedroom	7000.0	8.0	450.0	4	6428.571429
6711	Thanisandra	3 BHK	1806.0	6.0	116.0	3	6423.034330
8411	other	6 BHK	11338.0	9.0	1000.0	6	8819.897689

```
In [45]: df9= df8[df8.bath<df8.Bhk+2]
df9.shape
```

```
Out[45]: (7251, 7)
```

In [46]: df10= df9.drop(['size','price\_per\_sqft'],axis="columns")  
df10.head(3)

Out[46]:

	location	total_sqft	bath	price	Bhk
0	1st Block Jayanagar	2850.0	4.0	428.0	4
1	1st Block Jayanagar	1630.0	3.0	194.0	3
2	1st Block Jayanagar	1875.0	2.0	235.0	3

# Applying the model

In [47]: dummies = pd.get\_dummies(df10.location)  
dummies.head()

Out[47]:

	1st Block Jayanagar	1st Phase JP Nagar	2nd Phase Judicial Layout	2nd Stage Nagarbhavi	5th Block Hbr Layout	5th Phase JP Nagar	6th Phase JP Nagar	7th Phase JP Nagar	8th Phase JP Nagar	9th Phase JP Nagar	...	Vishveshwara Layout
0	1	0	0	0	0	0	0	0	0	0	...	
1	1	0	0	0	0	0	0	0	0	0	...	
2	1	0	0	0	0	0	0	0	0	0	...	
3	1	0	0	0	0	0	0	0	0	0	...	
4	1	0	0	0	0	0	0	0	0	0	...	

5 rows × 242 columns

In [48]: df11 = pd.concat([df10,dummies], axis = 'columns')

In [49]: df11.drop(['1st Block Jayanagar','location'],axis='columns')

Out[49]:

	total_sqft	bath	price	Bhk	1st Phase JP Nagar	2nd Phase Judicial Layout	2nd Stage Nagarbhavi	5th Block Hbr Layout	5th Phase JP Nagar	6th Phase JP Nagar	...	Vishveshwara Layout
0	2850.0	4.0	428.0	4	0	0	0	0	0	0	...	
1	1630.0	3.0	194.0	3	0	0	0	0	0	0	...	
2	1875.0	2.0	235.0	3	0	0	0	0	0	0	...	
3	1200.0	2.0	130.0	3	0	0	0	0	0	0	...	
4	1235.0	2.0	148.0	2	0	0	0	0	0	0	...	
...	...	...	...	...	...	...	...	...	...	...	...	
10232	1200.0	2.0	70.0	2	0	0	0	0	0	0	...	
10233	1800.0	1.0	200.0	1	0	0	0	0	0	0	...	
10236	1353.0	2.0	110.0	2	0	0	0	0	0	0	...	
10237	812.0	1.0	26.0	1	0	0	0	0	0	0	...	
10240	3600.0	5.0	400.0	4	0	0	0	0	0	0	...	

7251 rows × 245 columns

```
In [50]: df11.shape
```

```
Out[50]: (7251, 247)
```


```
In [51]: X = df11.drop(['price','location'], axis='columns')
```

```
In [52]: X.head()
```

```
Out[52]:
```

	total_sqft	bath	Bhk	1st Block Jayanagar	1st Phase JP Nagar	2nd Phase Judicial Layout	2nd Stage Nagarbhavi	5th Block Hbr Layout	5th Phase JP Nagar	6th Phase JP Nagar	...	Vishveshwa Lay
0	2850.0	4.0	4	1	0	0	0	0	0	0	...	
1	1630.0	3.0	3	1	0	0	0	0	0	0	...	
2	1875.0	2.0	3	1	0	0	0	0	0	0	...	
3	1200.0	2.0	3	1	0	0	0	0	0	0	...	
4	1235.0	2.0	2	1	0	0	0	0	0	0	...	

5 rows × 245 columns



```
In [53]: y= df11.price  
y.head()
```

```
Out[53]: 0    428.0  
1    194.0  
2    235.0  
3    130.0  
4    148.0  
Name: price, dtype: float64
```

```
In [54]: from sklearn.model_selection import train_test_split  
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.2,random_state=10)
```

```
In [55]: from sklearn.linear_model import LinearRegression  
lr = LinearRegression()  
lr.fit(X_train,y_train)  
lr.score(X_test,y_test)
```

```
Out[55]: 0.845227769787358
```

```
In [56]: from sklearn.neighbors import KNeighborsRegressor  
knn = KNeighborsRegressor(n_neighbors=3)  
knn.fit(X_train,y_train)  
knn.score(X_test,y_test)
```

```
Out[56]: 0.7211004574660052
```

```
In [57]: from sklearn.model_selection import ShuffleSplit  
from sklearn.model_selection import cross_val_score  
cv =ShuffleSplit(n_splits=5, test_size=0.2,random_state =0)  
cross_val_score(LinearRegression(),X,y,cv=cv)
```

```
Out[57]: array([0.82433489, 0.77166234, 0.85089567, 0.80836261, 0.83653286])
```

```

In [58]: from sklearn.model_selection import GridSearchCV

#from sklearn.linear_model import Lasso
from sklearn.tree import DecisionTreeRegressor

def find_best_model_using_gridsearchcv(X,y):
    algos = {
        'linear_regression' : {
            'model': LinearRegression(),
            'params': {
                'normalize': [True, False]
            }
        },
        'KNN_regressor':{
            'model': KNeighborsRegressor(),
            'params':{
                'weights':['uniform', 'distance'],
                'algorithm':['auto']
            }
        },
        'decision_tree': {
            'model': DecisionTreeRegressor(),
            'params': {
                'criterion' : ['mse','friedman_mse'],
                'splitter': ['best','random']
            }
        }
    }
    scores = []
    cv = ShuffleSplit(n_splits=5, test_size=0.2, random_state=0)
    for algo_name, config in algos.items():
        gs = GridSearchCV(config['model'], config['params'], cv=cv, return_train_score=False)
        gs.fit(X,y)
        scores.append({
            'model': algo_name,
            'best_score': gs.best_score_,
        })

    return pd.DataFrame(scores,columns=['model','best_score'])

find_best_model_using_gridsearchcv(X,y)

```

Out[58]:

	model	best_score
0	linear_regression	0.818362
1	KNN_regressor	0.714367
2	decision_tree	0.768092

```
In [59]: def predict_price(model,location,sqft,bath,bhk):  
        loc_index = np.where(X.columns==location)[0][0]  
  
        x = np.zeros(len(X.columns))  
        x[0] = sqft  
        x[1] = bath  
        x[2] = bhk  
        if loc_index >= 0:  
            x[loc_index] = 1  
  
        return model.predict([x])[0]
```

```
In [60]: predict_price(lr,'1st Phase JP Nagar',1000, 2, 2)
```

```
Out[60]: 83.49904676985425
```

```
In [61]: predict_price(knn,'1st Phase JP Nagar',1000, 2, 2)
```

```
Out[61]: 38.526666666666664
```

```
In [ ]:
```